

A Research Principle for the AI Era

Jianmin Liao

September 28, 2026

Suppose we allocate research resources over many cycles toward an ambitious goal: proving $P \neq NP$, creating artificial superintelligence, or turning science fiction into technologies that benefit humanity. The goal may take decades to achieve. Either only we believe it is possible, or it is useful only to us; in either case, others will not fund it. Progress toward this goal has intrinsic value to us, but consumes resources without bringing new ones in. Meanwhile, papers in leading venues and topics others will pay us to pursue can attract resources for the next cycle. We call this funding return external value. If the intrinsic goal is all we ultimately care about, how much should we allocate to it early in our careers?

Even when intrinsic value is the sole objective, the optimal initial allocation to it approaches zero as the horizon grows; with linear expected returns, it becomes exactly zero for sufficiently long horizons.

Goal

Maximize expected total intrinsic value over n research cycles:

$$\max \mathbb{E} \left[\sum_{i=1}^n I_i \right].$$

Initial intrinsic allocation approaches zero

The optimal initial intrinsic share approaches zero as the horizon grows, even with nonlinear returns.

Assumptions for Theorem 1

Consider n research cycles, starting with a seed budget $c_1 = 1$. In cycle i , expenditure is c_i and the fraction allocated to intrinsic work is $p_i \in [0, 1]$. The researcher observes the current budget and balance before choosing p_i .

1. Random return functions. Let $F_i(x)$ and $G_i(x)$ be the intrinsic and external returns from expenditure x , so

$$I_i = F_i(p_i c_i), \quad E_i = G_i((1 - p_i) c_i), \quad c_{i+1} = d_i E_i.$$

The nonnegative return functions vanish at zero. The pairs (F_i, G_i) are iid across cycles. The multipliers $d_i \geq 1$ are iid, bounded, and independent of the return functions. Allocations are chosen before the current draws. At every feasible state, finite-horizon optimal policies exist and have finite expected value.

2. A growing reserve. For every decision history, the lowest possible external return from expenditure x is bx , with the same b across cycles. The funding balance obeys

$$B_0 = 0, \quad B_i = B_{i-1} + E_i - c_i.$$

For every possible outcome, require

$$B_i \geq \rho B_{i-1}, \quad 1 < \rho < b.$$

Research income must cover spending while growing the reserve.

3. Positive return increments. Write $f(x) = \mathbb{E}[F_i(x)]$. For some $m > 0$ and every $x, h \geq 0$, require

$$f(x+h) - f(x) \geq mh, \quad G_i(x+h) - G_i(x) \geq bh.$$

The second inequality holds for every realization. These conditions allow nonlinear returns with a positive minimum gain from additional investment.

Approximate initial external focus

Theorem 1. *For every horizon $n \geq 2$, the first intrinsic share $p_1^{*(n)}$ under any optimal policy satisfies*

$$0 \leq p_1^{*(n)} \leq \frac{f(1)}{m(b-\rho)b^{n-2}}.$$

Consequently, $p_1^{(n)} \rightarrow 0$ as $n \rightarrow \infty$: the optimal initial allocation becomes arbitrarily close to full external focus.*

A nonlinear example

Use $n = 100$, $\rho = 1.01$, and

$$I_i = p_i c_i + A_i \sqrt{p_i c_i}, \quad E_i = 1.05(1 - p_i)c_i, \quad d_i = 1,$$

with iid $A_i \sim U[0.8, 1.2]$. Then $f(x) = x + \sqrt{x}$, so these assumptions hold with $m = 1$ and $b = 1.05$. We numerically maximize expected total intrinsic value and show one realization with random seed 20260928.




Figure 1: One 100-cycle realization of the nonlinear model. The solid line shows the intrinsic allocation share; the dashed line shows cumulative realized intrinsic value. The first share is 0.0074 percent, rising to about 3.8 percent in later cycles.

Proof of Theorem 1

Every feasible state satisfies $0 \leq B \leq c$: this holds initially and is preserved because

$$0 \leq B' = B + E - c \leq E \leq d_i E = c'.$$

Intrinsic expenditure t is feasible precisely when

$$0 \leq t \leq \frac{(b-1)c - (\rho-1)B}{b}.$$

Let $V_k(c, B)$ denote optimal expected intrinsic value with k cycles remaining. We first show that increasing the state by (δ, η) , with $0 \leq \eta \leq \delta$, gives

$$V_k(c + \delta, B + \eta) - V_k(c, B) \geq L_k \delta, \quad L_k = m(b - \rho)b^{k-2}.$$

For $k = 1$, optimal intrinsic expenditure is the upper endpoint above. It increases by at least $(1 - \rho/b)\delta$, so the lower bound on f gives $L_1 \delta$.

For the induction step, use the smaller state's optimal intrinsic expenditure in the larger state. External expenditure increases by δ ; the reserve constraint remains satisfied because $b\delta \geq \delta + (\rho - 1)\eta$. Couple the two decisions using the same return functions and funding multiplier. The external return increases by some $\Delta E \geq b\delta$. The next-state differences are

$$\delta' = d_i \Delta E, \quad \eta' = \eta + \Delta E - \delta,$$

so $0 \leq \eta' \leq \delta'$ and $\delta' \geq b\delta$. Current intrinsic returns are identical. The induction hypothesis therefore gives a continuation gain of at least $bL_{k-1}\delta = L_k\delta$.

Finally, at the initial state $(1, 0)$, compare an optimal intrinsic expenditure $t = p_1^{*(n)}$ with full external allocation. Full external allocation increases the next budget and balance by $(d_i \Delta E, \Delta E)$,

where $\Delta E \geq bt$. The value comparison just proved gives a continuation gain of at least $bL_{n-1}t = L_nt$. Optimality of t requires its current expected intrinsic return to cover this gain:

$$L_nt \leq f(t) \leq f(1).$$

Dividing by L_n proves the bound. Since $b > 1$, its right-hand side tends to zero. $\square$

Linear expected returns

This initial share becomes exactly zero for sufficiently long horizons with linear expected returns.

Assumptions for Theorem 2

Keep the same initial state and growing-reserve condition. Write $\mathbb{E}_i$ for expectation conditional on the information available when p_i is chosen, including that choice.

1. Linear expected returns. For every decision history and allocation,

$$\mathbb{E}_i[I_i] = ap_i c_i, \quad \mathbb{E}_i[E_i] = w(1 - p_i)c_i, \quad c_{i+1} = d_i E_i.$$

The constants $a > 0$ and w are common to all cycles. The multipliers $d_i \geq 1$ are iid and bounded; each draw is independent of the decision information and current external return.

2. Decision information. The researcher observes the current budget and balance and chooses p_i before the current returns and funding multiplier are drawn. Only b , w , and $\mathbb{E}[d_i]$ are known about their distributions.

Full initial external focus

Theorem 2. *There is a finite integer $K \geq 1$, independent of the horizon n , such that whenever $n > K$, the unique optimal decision in each of the first $n - K$ cycles is to allocate all resources to external value:*

$$p_i = 0 \quad (1 \leq i \leq n - K).$$

A numerical example




Figure 2: One optimal trajectory with iid random coefficients. All resources go to external value in the first eight cycles. The line shows the intrinsic allocation share. Each bar shows cumulative intrinsic value, with the current increment highlighted.

Use $I_i = a_i p_i c_i$ and $E_i = b_i(1 - p_i)c_i$, with $n = 20$, $\rho = 1.045$, and independent uniform draws

$$a_i \sim U[1.5, 2.5], \quad b_i \sim U[1.05, 1.15], \quad d_i \sim U[1, 1.2].$$

Here $K = 12$, so the first eight cycles are devoted entirely to external value. The figure shows one realization with random seed 20260928.

Proof of Theorem 2

Every feasible state satisfies $0 \leq B \leq c$, as shown above. Write $u = (1 - p)c$ for external expenditure. The reserve constraint gives the feasible interval

$$\frac{c + (\rho - 1)B}{b} \leq u \leq c.$$

Divide the objective by a , giving expected stage reward $c - u$. Independence of the funding multiplier gives

$$v = w\mathbb{E}[d_i] \geq w \geq b > \rho > 1, \quad \mathbb{E}_i[c_{i+1}] = v(1 - p_i)c_i.$$

At every decision history with state (c, B) and m cycles remaining, the value is $V_m(c, B) = A_m c + D_m B$, with $A_0 = D_0 = 0$. To see this, define $S_m = vA_m + wD_m$. The Bellman objective is

$$(1 - D_{m-1})c + D_{m-1}B + (S_{m-1} - 1)u.$$

Its maximizing endpoint depends only on S_{m-1} , and both endpoints are linear in (c, B) , completing the induction.

When $S_{m-1} \leq 1$, choose the lower endpoint. Starting from $S_0 = D_0 = 0$, this gives

$$D_m = D_{m-1} + \frac{(\rho - 1)(S_{m-1} - 1)}{b} \leq 0$$

and

$$\begin{aligned} S_m &= \frac{v + w(\rho - 1)}{b} S_{m-1} + \frac{v(b - 1) - w(\rho - 1)}{b} \\ &\quad + (w - v)D_{m-1} \geq \rho S_{m-1} + b - \rho. \end{aligned}$$

Since $\rho > 1$ and $b > \rho$, this sequence crosses 1 in finite time. Let $K = \min\{m \geq 1 : S_m > 1\}$.

Once $S_{m-1} > 1$, the unique maximizing choice is $u = c$. Then $D_m = D_{m-1} \leq 0$ and

$$S_m = vS_{m-1} + (w - v)D_{m-1} > S_{m-1} > 1.$$

Thus every decision with more than K cycles remaining uniquely allocates everything to external value. These are precisely the first $n - K$ cycles. $\square$

Even when only intrinsic value matters, the optimal starting allocation is entirely or almost entirely external.